



[PYTHON \(https://www.analyticsvidhya.com/blog/category/python-2/\)](https://www.analyticsvidhya.com/blog/category/python-2/)

[REINFORCEMENT LEARNING \(https://www.analyticsvidhya.com/blog/category/machine-learning/reinforcement-learning/\)](https://www.analyticsvidhya.com/blog/category/machine-learning/reinforcement-learning/)

## Nuts and Bolts of Reinforcement Learning: Introduction to Temporal Difference (TD) Learning


ANKIT CHOUDHARY (<https://www.analyticsvidhya.com/blog/author/ankit2106/>), MARCH 28, 2019 [LOGIN TO BOOKMARK THIS ARTICLE \(https://id.analyticsvidhya.com/accounts/...](https://id.analyticsvidhya.com/accounts/...)

### Introduction

Q-learning became a household name in [data science \(https://courses.analyticsvidhya.com/courses/introduction-to-data-science-2?utm\\_source=blog&utm\\_medium=reinforcement-learning-temporal-difference\)](https://courses.analyticsvidhya.com/courses/introduction-to-data-science-2?utm_source=blog&utm_medium=reinforcement-learning-temporal-difference) when DeepMind came up with an algorithm that reached superhuman levels on ATARI games. It's one of the core components of reinforcement learning (RL). I regularly come across Q-learning whenever I'm reading up about RL.

But what does Q-learning have to do with our topic of temporal difference learning? Let me take an example to give an intuition of what temporal difference learning is about.

Rajesh is planning to travel to Jaipur from Delhi in his car. A quick check on Google Maps shows him an estimated 5 hour journey. Unfortunately, there is an unexpected delay due to a roadblock (anyone who has taken a long journey can relate to this!). The estimated arrival time for Rajesh now jumps up to 5 hours 30 minutes.



Halfway through the journey, he finds a bypass that reduces his arrival time. So overall, his journey from Delhi to Jaipur takes 5 hours 10 minutes.

Email Id

Contact Number



Download Resource



Did you notice how Rajesh's arrival time kept changing and updating based on different episodes? This, in a nutshell, illustrates the temporal difference learning concept. In this article, I will introduce you to this algorithm and its components in detail, including how Q-learning fits into the picture. We'll also pick up a case study and solve it in Python.

This is one of the most intriguing reinforcement learning concepts so let's have fun learning this!

## Table of Contents

- Introduction to Temporal Difference (TD) Learning
  - TD Prediction
  - TD Control
- Off-Policy vs On-Policy Learning
- Q-learning: Off-Policy Temporal Difference Learning
- SARSA: On-Policy Temporal Difference Learning
- Case Study: Taxi Scheduling using Q-learning in Python

## Introduction to Temporal Difference (TD) Learning

We developed an intuition for temporal difference learning in the introduction. Now let's understand it in a bit more detail.

In my [previous article on Monte Carlo Learning](https://www.analyticsvidhya.com/blog/2018/11/reinforcement-learning-introduction-monte-carlo-learning-openai-gym/?utm_source=blog&utm_medium=reinforcement-learning-temporal-difference) ([https://www.analyticsvidhya.com/blog/2018/11/reinforcement-learning-introduction-monte-carlo-learning-openai-gym/?utm\\_source=blog&utm\\_medium=reinforcement-learning-temporal-difference](https://www.analyticsvidhya.com/blog/2018/11/reinforcement-learning-introduction-monte-carlo-learning-openai-gym/?utm_source=blog&utm_medium=reinforcement-learning-temporal-difference)), we learned how to use it for solving the Markov Decision Process (MDP) when the model dynamics of the environment are not known in advance. So why don't we just use that instead of TD?

Well, although Monte Carlo learning provides an efficient and simple approach for model-free learning, there are some limitations to it. It can be applied only for episodic tasks.

An episodic task lasts a finite amount of time. For example, playing a single game of Chess is an episodic task, which you win or lose. If an episode is very long (which it is in many cases), then we have to wait a long time to compute value functions.

Temporal difference (TD) learning, which is a model-free learning algorithm, has two important properties:

- It doesn't require the model dynamics to be known in advance

**Your Ultimate path for Becoming a  
DATA Scientist!**

Download this learning path to start your data science journey.


Name

Email Id

Contact Number

 [Download Resource](#)

- It can be applied for non-episodic tasks as well



The TD learning algorithm was introduced by the great Richard Sutton in 1988. The algorithm takes the benefits of both the Monte Carlo method and dynamic programming (DP) into account:

- Like the Monte Carlo method, it doesn't require model dynamics, and
- Like dynamic programming, it doesn't need to wait until the end of the episode to make an estimate of the value function

Instead, temporal difference learning approximates the current estimate based on the previously learned estimate. This approach is also called bootstrapping.

## Getting the Intuition Behind TD Prediction

We try to predict the state values in temporal difference learning, much like we do in Monte Carlo prediction and dynamic programming prediction. In Monte Carlo prediction, we estimate the value function by simply taking the mean return for each state whereas in Dynamic Programming and TD Learning, we update the value of a previous state by the current state. But TD learning does not need model of the environment unlike DP.

How can we do this? TD learning uses something called a TD update rule for updating the value of a state:

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

value of a previous state = value of previous state + learning\_rate \* (reward + discount\_factor(value of current state) - value of previous state)

What does this equation actually mean?

## Your Ultimate path for Becoming a DATA Scientist

Don't lose the value function by simply taking the mean return for each state

Name

Email Id

Contact Number

Download Resource

This is the difference between the **actual reward** ( $r + \text{Gamma} * V(s')$ ) and the expected reward  $V(s)$  multiplied by the learning rate  $\alpha$ .

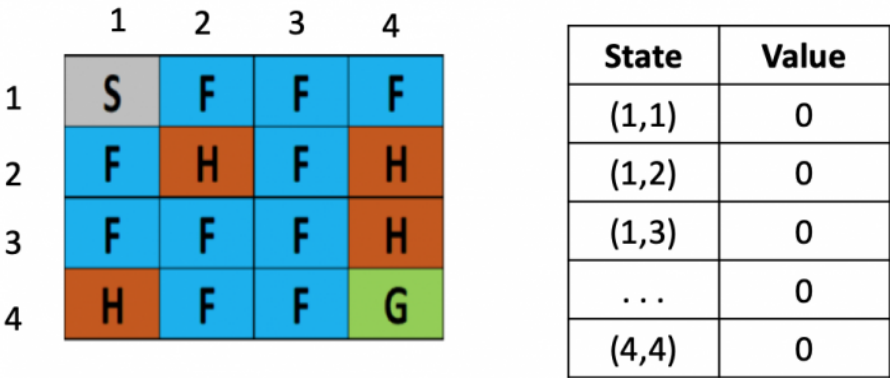
*What does the learning rate signify?*

The learning rate, also called step size, is useful for convergence.

Since we take the difference between the actual and predicted values, this is like an error. We can call it a **TD error**. Notice that the TD error at each time is the error in the estimate made at that time. Because the TD error depends on the next state and next reward, it is not actually available until one timestep later. Iteratively, we will try to minimize this error.

### Understanding TD Prediction using the Frozen Lake Example

Let us understand TD prediction with the frozen lake example. The frozen lake environment is shown next. First, we will initialize the value function as 0, as in  $V(s)$  as 0 for all states, as shown in the following state-value diagram:



Say we are in a starting state ( $s$ ) (1,1) and we take an action right and move to the next state ( $s'$ ) (1,2) and receive a reward ( $r$ ) as -0.4.

How can we update the value of the state using this information? Recall the TD update equation:

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Your Ultimate path for Becoming a DATA Scientist!

Download this learning path to start your data science journey.

Let us consider the learning rate ( $\alpha$ ) as 0.1 and the discount factor ( $\gamma$ ) as 0.5; we know that the value of the state (1,1), as in  $v(s)$ , is 0 and the value of the next state (1,2), as in  $V(s')$ , is also 0. The reward ( $r$ ) we obtained is -0.3. We substitute this in the TD rule as follows:

$V(s) = 0 + 0.1 [ -0.4 + 0.5 (0)-0]$

$V(s) = - 0.04$

So, we update the value for the state (1,1) as -0.04 in the value table, as shown in the following diagram:

Name

Email Id

Contact Number

Download Resource

	Right			
	1	2	3	4
1	S → F	F	F	F
2	F	H	F	H
3	F	F	F	H
4	H	F	F	G

State	Value
(1,1)	-0.04
(1,2)	0
(1,3)	0
...	0
(4,4)	0

Now that we are in the state (s) as (1,2), we take an action right and move to the next state (s') (1,3) and receive a reward (r) -0.4. How do we update the value of the state (1, 2) now? We will substitute the values in the TD update equation:

$$V(s) = 0 + 0.1 [-0.4 + 0.5(0) - 0]$$

$$V(s) = -0.04$$

Go ahead and update the value of state (1,2) as -0.04 in the value table:

		Right			
		1	2	3	4
1	S	F	→ F	F	F
2	F	H	F	H	
3	F	F	F	H	
4	H	F	F	G	

State	Value
(1,1)	-0.04
(1,2)	-0.04
(1,3)	0
...	0
(4,4)	0

With me so far? We are now in the state (s) (1,3). Let's take an action left.

We again go back to that state (s') (1,2) and we receive a reward (r) -0.4. The value of the state (1,3) is 0 and the value of the next state (1,2) is -0.04 in the value table. Now we can update the value of state (1,3) as follows:

$$V(s) = 0 + 0.1 [-0.4 + 0.5 (-0.04) - 0]$$

$$V(s) = 0.1[-0.42]$$

$$V(s) = -0.042$$

You know what to do now. Update the value of state (1,3) as -0.042 in the value table:

## Your Ultimate path for Becoming a DATA Scientist!

Do you want to learn the first step to start your data science journey.





Download Resource

		Left			
		1	2	3	4
1	S	F	F ←	F	F
2	F	H	F	H	
3	F	F	F	H	
4	H	F	F	G	

State	Value
(1,1)	-0.04
(1,2)	-0.04
(1,3)	-0.042
...	0
(4,4)	0

We update the value of all the states in a similar fashion using the TD update rule. To summarize, the steps involved in the TD prediction algorithm are:

1. First, initialize  $V(S)$  to 0 or some arbitrary value
2. Then, begin the episode. For every step in the episode, perform an action  $A$  in the state  $S$  and receive a reward  $R$  and move to the next state ( $s'$ )
3. Update the value of the previous state using the TD update rule
4. Repeat steps 2 and 3 until we reach the terminal state

## Understanding Temporal Differencing Control

In Temporal Difference prediction, we *estimated* the value function. In TD control, we *optimize* the value function. There are two kinds of algorithms we use for TD Control:

- Off-policy learning algorithm: Q-learning
- On-policy learning algorithm: SARSA

## Off-Policy vs On-Policy

What's the difference between off-policy and on-policy learning? The answer lies in their names:

- **Off-policy learning:** The agent learns about policy  $\pi$  from experience sampled from another policy  $\mu$
- **On-policy learning:** The agent learns about policy  $\pi$  from experience sampled from the same policy  $\pi$

Let me break this down in the form of an example. Let's say you joined a new firm as a data scientist. In this scenario, you can equate on-policy learning as learning on the job. You'll be trying different things and learning only from your own experience.

Off-policy learning would be where you have full access to the actions of another employee. All you do in this scenario is learn from that employee's experience and not repeat something that the employee has failed at.

## Q-learning

## Your Ultimate path for Becoming a DATA Scientist!

Download this learning path to start your data science journey.

Name

Email Id

Contact Number



Download Resource

Q-learning is a very popular and widely used off-policy TD control algorithm.

In Q learning, our concern is the state-action value pair—the effect of performing an action a in the state s. This tells us how good an action is for the agent at a particular state (Q(s,a)), rather than looking only at how good it is to be in that state (V(s))

We will update the Q value based on the following equation:


$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Why is Q-learning considered as an off-policy technique? This is because it updates its Q-values using the Q-value of the next state s' and the greedy action a'. In other words, it estimates the return (total discounted future reward) for state-action pairs assuming a greedy policy (maxQ(s'a')) was followed despite the fact that it's not following a greedy policy!

The above equation is similar to the TD prediction update rule with a subtle difference. Below are the steps involved in Q-learning (I want you to notice the difference here):

- 1. First, initialize the Q function to some arbitrary value
- 2. Take an action from a state using epsilon-greedy policy () and move it to the new state
- 3. Update the Q value of a previous state by following the update rule
- 4. Repeat steps 2 and 3 till we reach the terminal state

Now, let's go back to our example of Frozen Lake. Let us say we are in a state (3,2) and have two actions (left and right). Refer to the below figure:




State	Action	Value
(3,2)	Left	0.2
(3,2)	Right	0.4

Your Ultimate path for Becoming a DATA Scientist!

Download this learning path to start your data science journey.

We select an action using the epsilon-greedy policy in Q-learning. We either explore a new action with the probability epsilon or we select the best action with a probability 1 – epsilon. Suppose we select a probability epsilon and select a particular action (moving down):

 Download Resource



	1	2	3	4
1	S	F	F	F
2	F	H	F	H
3	F	F	F	H
4	H	F	F	G

State	Action	Value
(3,2)	Left	0.2
(3,2)	Right	0.4
(3,2)	Down	0.6
(4,2)	Up	0.2
(4,2)	Down	0.4
(4,2)	Right	0.6

We have performed a downward action in the state (3,2) and reached a new state (4,2) using the epsilon-greedy policy. How do we update the value of the previous state (3,2) using our update rule? It's pretty straightforward!

Let us consider alpha as 0.1, the discount factor as 1 and reward as 0.4:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

$$Q(3,2 \text{ down}) = Q(3,2 \text{ down}) + 0.1 (0.4 + 1 \max [Q(4,2 \text{ action})] - Q(3,2 \text{ down}))$$

We can say the value of a state (3,2) with downward action is 0.6 in the Q table. What is max Q ( (4,2), action) for the state (4,2)?

We have explored three actions (up, down, and right) so we will take the maximum value based on these actions only. There is no exploration involved here – this is a straightforward greedy policy.

Based on the previous Q table, we can plug in the values:

$$Q(3,2, \text{down}) = 0.6 + 0.1 (0.4 + 1 * \max [0.2, 0.4, 0.6] - 0.6)$$

$$Q(3,2, \text{down}) = 0.64$$

So, we update the value of Q ((3,2), down) to 0.64.

Now, we are in the (4,2) state. What action should we perform? Based on the epsilon-greedy policy, we can either explore a new action with a probability epsilon, or select the best action with a probability 1-epsilon. Let us say we select the latter option. So, in (4,2), the action right has a maximum value and that's what we'll select:


	1	2	3	4
1	S	F	F	F
2	F	H	F	H
3	F	F	F	H
4	H	F	F	G

## Your Ultimate path for Becoming a

### DATA Scientist

Don't waste your time. Start your data science journey.

State	Action	Value
(4,2)	Up	0.2
(4,2)	Down	0.4
(4,2)	Right	0.6

 Download Resource



Right, we are have moved to the state (4,3). Things are shaping up nicely so far. But wait – how do we update the value of the previous state?

$$Q( (4,2), right) = Q( (4,2), right ) + 0.1 ( 0.4 + 1 \cdot \max [Q( (4,3) \text{ action} ) ] - Q( (4,2), right)$$

If you look at the Q table that follows, we have explored only two actions (up and down) for the state (4,3). So, we will take a maximum value based only on these actions (we will not perform an epsilon-greedy policy here; we simply select the action which has maximum value):

	1	2	3	4
1	S	F	F	F
2	F	H	F	H
3	F	F	F	H
4	H	F	F	G

State	Action	Value
(4,2)	Up	0.2
(4,2)	Down	0.4
(4,2)	Right	0.6
(4,3)	Up	0.2
(4,3)	Down	0.4

$$Q ( (4,2), right) = Q((4,2),right) + 0.1 (0.4 + 1 \max [ (Q (4,3), up) , ( Q(4,3),down) ] - Q ((4,2), right)$$

$$\begin{aligned}
 Q ( (4,2), right) &= 0.6 + 0.1 (0.4 + 1 \max [ 0.2,0.4] - 0.8) \\
 &= 0.6 + 0.1 (0.4 + 1(0.4) - 0.6) \\
 &= 0.62
 \end{aligned}$$

Awesome! We'll update the value of the state Q((4,2), right) to 0.62.

*This is how we get the state-action values in Q-learning. We use the epsilon-greedy policy to decide what action to take, and simply pick the maximum action while updating the Q value.*

Your Ultimate path for Becoming a  
DATA Scientist!

Download this learning path to start your data science journey.

### SARSA

State-Action-Reward-State-Action (SARSA) is an on-policy TD control algorithm. Similar to what we did in Q-learning, we **focus on state-action value instead of a state-value pair**. In SARSA, we update the Q value based on the below update rule:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

You might have noticed that **there is no max Q(s',a')** (unlike Q-learning). Here, it is simply Q(s,a). You'll understand this when you go through the below SARSA steps:

Name

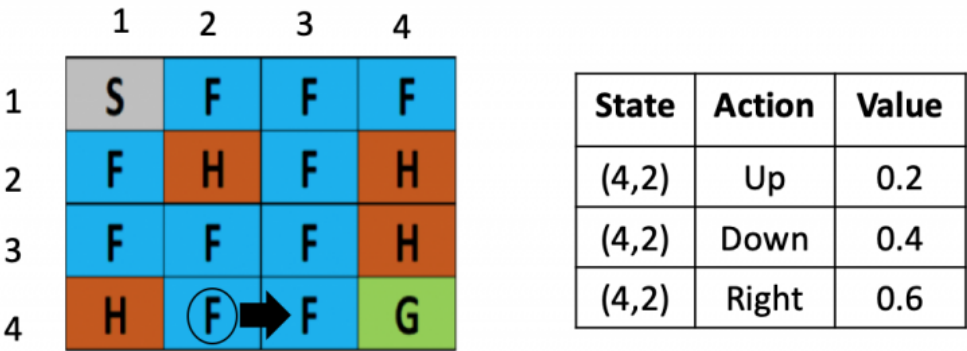
Email Id

Contact Number

Download Resource

1. First, initialize the Q values to some arbitrary values
2. Select an action by the epsilon-greedy policy () and move from one state to another
3. Update the Q value in the previous state by following the update rule, where a' is the action selected by an epsilon-greedy policy ()
4. Now we'll understand the algorithm step-by-step

Let us consider the same frozen lake example. Suppose we are in state (4,2). We decide the action based on the epsilon-greedy policy. Let's say we use a probability 1 – epsilon and select the best action (moving right):

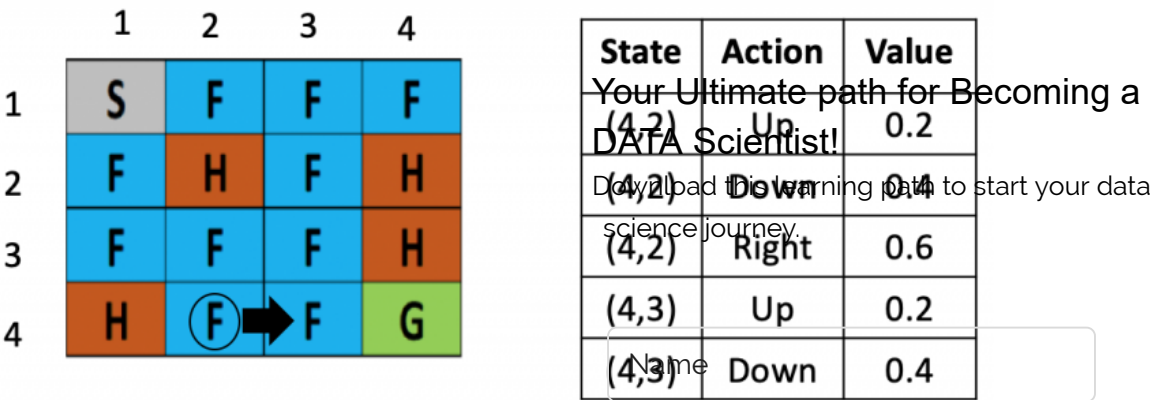


We land in the state (4,3). How do we update the value of the previous state (4,2)? I'll pull up the above equation and plug in the values. Let us consider the alpha as 0.1, the reward as 0.4, and discount factor as 1:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$

$$Q( (4,2), \text{right} ) = Q( (4,2), \text{right} ) + 0.1 ( 0.4 + 1 * Q( (4,3), \text{action} ) ) - Q( (4,2), \text{right} )$$

How do we choose the value for  $Q( (4,3), \text{action} )$ ? We can't just pick up  $\max ( Q( (4,3), \text{action} ) )$  like we did in Q-learning. In SARSA, we use the epsilon-greedy policy. Look at the Q table I've shown below. In state (4,3), we have explored two actions:



We either explore with a probability epsilon or exploit with a probability 1 – epsilon. Let's say we select the former option and explore a new action (moving right):

Email Id

Contact Number


Download Resource

	1	2	3	4
1	S	F	F	F
2	F	H	F	H
3	F	F	F	H
4	H	F	F	G

State	Action	Value
(4,2)	Up	0.2
(4,2)	Down	0.4
(4,2)	Right	0.6
(4,3)	Up	0.2
(4,3)	Down	0.4
(4,3)	Right	0.7

$$Q((4,2), \text{right}) = Q((4,2), \text{right}) + 0.1 (0.4 + 1 (Q((4,3), \text{right}) - Q((4,2), \text{right})))$$

$$Q((4,2), \text{right}) = 0.6 + 0.1 (0.4 + 1(0.7) - 0.6)$$

$$Q((4,2), \text{right}) = 0.65$$

*And that's how we get the state-action values in SARSA. We take the action using the epsilon-greedy policy and pick the action using the epsilon-greedy policy updating the Q value.*


## Case Study in Python: Taxi Scheduling using Q-learning

We've finally come to the implementation part of our article! Time to fire up our Jupyter notebook and get our hands dirty. We will demonstrate Q-learning in action with an OpenAI Gym environment called Taxi-V2 (<https://gym.openai.com/envs/Taxi-v2/>).

*The aim of this project – our agent must **pick up the passenger at one location and drop him off at the destination/goal as fast as possible.***

**Your Ultimate path for Becoming a DATA Scientist!**

Download this learning path to start your data science journey.






R, G, Y and B in the above image are the 4 locations we will consider for our project. A few pointers:

- You receive +20 points for a successful drop-off



Download Resource

- Lose 1 point for every timestep the agent takes
- There is also a 10-point penalty for illegal pick-up and drop-off actions (if you don't drop the passenger in one of the 3 other locations)

Let's begin!

## Step 1: Import dependencies & Create the Environment

```
import numpy as np

import gym

import random

env = gym.make("Taxi-v2")

env.render()
```

## Step 2: Create the Q-table and initialize it

- We'll create our Q-table to know how many rows (states) and columns (actions) we need. We need to calculate the *action\_size* and the *state\_size*
- OpenAI Gym provides us a way to do that: `env.action_space.n` and `env.observation_space.n`

```
action_size = env.action_space.n
print("Action size ", action_size)

state_size = env.observation_space.n
print("State size ", state_size)
```

```
Action size  6
State size  500
```

```
qtable = np.zeros((state_size, action_size))
print(qtable)
```

```
[[0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0.]
 ...
 [0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0.]
```

## Your Ultimate path for Becoming a DATA Scientist!

Download this learning path to start your data science journey.




 [Download Resource](#)

## Step 3: Setting Hyperparameters

We'll specify the hyperparameters at this step:

```
total_episodes = 50000 # Total episodes
total_test_episodes = 100 # Total test episodes
max_steps = 99 # Max steps per episode

learning_rate = 0.7 # Learning rate
gamma = 0.618 # Discounting rate

# Exploration parameters
epsilon = 1.0 # Exploration rate
max_epsilon = 1.0 # Exploration probability at start
min_epsilon = 0.01 # Minimum exploration probability
decay_rate = 0.01 # Exponential decay rate for exploration prob
```

## Step 4: The Q-learning Algorithm

Time to put our theoretical learning to the test! Let's implement the Q-learning algorithm in Python:

## Your Ultimate path for Becoming a DATA Scientist!

Download this learning path to start your data science journey.



Download Resource

```

# 2 For life or until learning is stopped
for episode in range(total_episodes):
    # Reset the environment
    state = env.reset()
    step = 0
    done = False

    for step in range(max_steps):
        # 3. Choose an action a in the current world state (s)
        ## First we randomize a number
        exp_exp_tradeoff = random.uniform(0,1)

        ## If this number > greater than epsilon --> exploitation (taking the biggest Q value for this state)
        if exp_exp_tradeoff > epsilon:
            action = np.argmax(qtable[state,:])

        # Else doing a random choice --> exploration
        else:
            action = env.action_space.sample()

        # Take the action (a) and observe the outcome state(s') and reward (r)
        new_state, reward, done, info = env.step(action)

        # Update Q(s,a):= Q(s,a) + lr [R(s,a) + gamma * max Q(s',a') - Q(s,a)]
        qtable[state, action] = qtable[state, action] + learning_rate * (reward + gamma *
        np.max(qtable[new_state, :]) - qtable[state, action])

        # Our new state is state
        state = new_state

        # If done : finish episode
        if done == True:
            break

        # Reduce epsilon (because we need less and less exploration)
        epsilon = min_epsilon + (max_epsilon - min_epsilon)*np.exp(-decay_rate*episode)

```

## Your Ultimate path for Becoming a DATA Scientist!

Download this learning path to start your data science journey.





Download Resource

### Step 5: Use our Q-table to play Taxi!

After 50,000 episodes, our Q-table can be used as a "cheatsheet" to play the role of the Taxi.

```

env.reset()

rewards = []

for episode in range(total_test_episodes):
    state = env.reset()
    step = 0
    done = False
    total_rewards = 0
    #print("*****")
    #print("EPISODE ", episode)

    for step in range(max_steps):
        # UNCOMMENT IT IF YOU WANT TO SEE OUR AGENT PLAYING
        # env.render()

        # Take the action (index) that have the maximum expected future reward given that state
        action = np.argmax(qtable[state,:])

        new_state, reward, done, info = env.step(action)

        total_rewards += reward

    if done:
        rewards.append(total_rewards)
        #print ("Score", total_rewards)
        break

    state = new_state
env.close()

print ("Score over time: " + str(sum(rewards)/total_test_episodes))

```

Score over time: 8.4

## Your Ultimate path for Becoming a DATA Scientist!

As you can see the overall score is 8.4. This basically says that on an average the agent is scoring 8.4 points each episode after learning from 50000 episodes. Try increasing the maximum episodes and you will see that the score on the test episodes will increase.

## End Notes

TD Learning is the most widely used reinforcement learning method today. This is probably due to their great simplicity: they can be applied online, with a minimal amount of computation, to experience generated from interaction with an environment; they can be expressed nearly completely by single equations that can be implemented with small computer programs.





Download Resource



There are extensions to TD learning which makes them slightly more complicated and significantly more powerful. The special cases of TD methods introduced in the article should rightly be called one-step, tabular, model-free TD methods. We can extend them to n-step forms (a link to Monte Carlo methods). We can also extend them to various forms of function approximation rather than tables (a link to deep learning and artificial neural networks) which is where Deep Q Learning comes into the picture.

Here's a challenge for the readers. I encourage everyone to use the provided implementation to implement SARSA algorithm. Let me know how it works out and share your code in the comment section.

You can also read this article on Analytics Vidhya's Android APP



[https://play.google.com/store/apps/details?id=com.analyticsvidhya.android&utm\\_source=blog\\_article&utm\\_campaign=blog&pcampaignid=MKT-Other-global-all-co-prtnr-py-PartBadge-Mar2515-1](https://play.google.com/store/apps/details?id=com.analyticsvidhya.android&utm_source=blog_article&utm_campaign=blog&pcampaignid=MKT-Other-global-all-co-prtnr-py-PartBadge-Mar2515-1)


#### Share this:

- (<https://www.analyticsvidhya.com/blog/2019/03/reinforcement-learning-temporal-difference-learning/?share=linkedin&nb=1>)
- (<https://www.analyticsvidhya.com/blog/2019/03/reinforcement-learning-temporal-difference-learning/?share=facebook&nb=1>)
- (<https://www.analyticsvidhya.com/blog/2019/03/reinforcement-learning-temporal-difference-learning/?share=twitter&nb=1>)
- (<https://www.analyticsvidhya.com/blog/2019/03/reinforcement-learning-temporal-difference-learning/?share=pocket&nb=1>)
- (<https://www.analyticsvidhya.com/blog/2019/03/reinforcement-learning-temporal-difference-learning/?share=reddit&nb=1>)

#### Related Articles



<https://www.analyticsvidhya.com/blog/2019/01/datahack-radio-reinforcement-learning-xander/>  
DataHack Radio #15: Exploring the Applications & Potential of Reinforcement Learning with Xander Steenbrugge  
(<https://www.analyticsvidhya.com/blog/2019/01/datahack-radio-reinforcement-learning-xander/>)  
January 6, 2019  
In "Podcast"



<https://www.analyticsvidhya.com/blog/2019/04/introduction-deep-q-learning-python/>  
A Hands-On Introduction to Deep Q-Learning using OpenAI Gym in Python  
(<https://www.analyticsvidhya.com/blog/2019/04/introduction-deep-q-learning-python/>)  
April 18, 2019  
In "Python"



<https://www.analyticsvidhya.com/blog/2017/01/introduction-to-reinforcement-learning-implementation/>  
Download this learning path to start your data science journey.  
Simple Beginner's guide to Reinforcement Learning & its implementation  
(<https://www.analyticsvidhya.com/blog/2017/01/introduction-to-reinforcement-learning-implementation/>)  
January 19, 2017  
In "Machine Learning"

Download Resource

TAGS : [PYTHON \(https://www.analyticsvidhya.com/blog/tag/python/\)](https://www.analyticsvidhya.com/blog/tag/python/), [REINFORCEMENT LEARNING \(https://www.analyticsvidhya.com/blog/tag/reinforcement-learning/\)](https://www.analyticsvidhya.com/blog/tag/reinforcement-learning/)

NEXT ARTICLE

## Computer Vision Tutorial: A Step-by-Step Introduction to Image Segmentation Techniques (Part 1)

(<https://www.analyticsvidhya.com/blog/2019/04/introduction-image-segmentation-techniques-python/>)

...

PREVIOUS ARTICLE

## 16 OpenCV Functions to Start your Computer Vision journey (with Python code)

(<https://www.analyticsvidhya.com/blog/2019/03/opencv-functions-computer-vision-python/>)



(<https://www.analyticsvidhya.com/blog/author/ankit2106/>)

### Ankit Choudhary (https://www.analyticsvidhya.com/blog/author/ankit2106/)

IIT Bombay Graduate with a Masters and Bachelors in Electrical Engineering. I have previously worked as a lead decision scientist for Indian National Congress deploying statistical models (Segmentation, K-Nearest Neighbours) to help party leadership/Team make data-driven decisions. My interest lies in putting data in heart of business for data-driven decision making.

[in](https://www.linkedin.com/in/ankit-choudhary-b9360826/) (<https://www.linkedin.com/in/ankit-choudhary-b9360826/>)

## 2 COMMENTS



NEHA AGARWAL

April 1, 2019 at 1:31 am (<https://www.analyticsvidhya.com/blog/2019/03/reinforcement-learning-temporal-difference-learning/#comment-157646>)

Excellent information for easily understandable even for common machine learning journey.

**Your Ultimate path for Becoming a DATA Scientist!**

Download this learning path to start your data science journey.

[Reply](#)



SADIA SIDDIQUI

May 5, 2019 at 5:05 pm (<https://www.analyticsvidhya.com/blog/2019/03/reinforcement-learning-temporal-difference-learning/#comment-158103>)

Very well explained. Thank you.

Name

[Reply](#)

Email Id

Contact Number



Download Resource

LEAVE A REPLY

Your email address will not be published.

Comment

Name (required)

Email (required)

Website

SUBMIT COMMENT


☐ Notify me of new posts by email.

JOIN THE NEXTGEN DATA SCIENCE ECOSYSTEM

- Get access to free courses on Analytics Vidhya
- Get free downloadable resource from Analytics Vidhya
- Save your articles
- Participate in hackathons and win prizes

([https://id.analyticsvidhya.com/accounts/login/?next=https://www.analyticsvidhya.com/blog/?utm\\_source=blog-subscribe&utm\\_medium=web](https://id.analyticsvidhya.com/accounts/login/?next=https://www.analyticsvidhya.com/blog/?utm_source=blog-subscribe&utm_medium=web))

Your Ultimate path for Becoming a DATA Scientist!



Download this learning path to start your data science journey.

Name

Email Id

Contact Number

Download Resource



([https://software.seek.intel.com/DataCenter\\_to\\_Edge\\_REG?](https://software.seek.intel.com/DataCenter_to_Edge_REG?registration_source=AnalyticsVidhya-APJ)

registration\_source=AnalyticsVidhya-APJ)

## POPULAR POSTS

24 Ultimate Data Science Projects To Boost Your Knowledge and Skills (& can be accessed freely)

(<https://www.analyticsvidhya.com/blog/2018/05/24-ultimate-data-science-projects-to-boost-your-knowledge-and-skills/>)

Essentials of Machine Learning Algorithms (with Python and R Codes)

(<https://www.analyticsvidhya.com/blog/2017/09/common-machine-learning-algorithms/>)

7 Types of Regression Techniques you should know!

(<https://www.analyticsvidhya.com/blog/2015/08/comprehensive-guide-regression/>)

A Complete Tutorial to Learn Data Science with Python from Scratch

(<https://www.analyticsvidhya.com/blog/2016/01/complete-tutorial-learn-data-science-python-scratch-2/>)

Understanding Support Vector Machine algorithm from examples (along with code)

(<https://www.analyticsvidhya.com/blog/2017/09/understaing-support-vector-machine-example-code/>)

Stock Prices Prediction Using Machine Learning and Deep Learning Techniques (with Python codes)

(<https://www.analyticsvidhya.com/blog/2018/10/predicting-stock-prices-using-machine-learning-and-deep-learning-techniques-python/>)

Introduction to k-Nearest Neighbors: Simplified (with implementation in Python)

(<https://www.analyticsvidhya.com/blog/2018/03/introduction-k-neighbours-algorithm-clustering/>)

A Simple Introduction to ANOVA (with applications in Excel)

(<https://www.analyticsvidhya.com/blog/2018/01/anova-analysis-of-variance/>)

**Your Ultimate path for Becoming a  
DATA Scientist!**

Download this learning path to start your data science journey

## RECENT POSTS

Top 7 Machine Learning Github Repositories for Data Scientists (<https://www.analyticsvidhya.com/blog/2019/06/top-7-machine-learning-github-repositories-data-scientists/>)

JUNE 6, 2019

The AI Comic: Z.A.I.N – Issue #1: Automating Attendance using Computer Vision (<https://www.analyticsvidhya.com/blog/2019/06/ai-comic-zain-issue-1-automating-computer-vision/>)

JUNE 3, 2019

DataHack Radio #23: Ines Montani and Matthew Honnibal – The Brains behind spaCy (<https://www.analyticsvidhya.com/blog/2019/06/datahack-radio-ines-montani-matthew-honnibal-brains-behind-spacy/>)

JUNE 3, 2019


Exclusive Interview with Sonny Laskar – Kaggle Master and Analytics Vidhya Hackathon Expert (<https://www.analyticsvidhya.com/blog/2019/05/exclusive-interview-sonny-laskar-kaggle-master-analytics-vidhya-hackathon-expert/>)

MAY 30, 2019



([https://datahack.analyticsvidhya.com/contest/intel-ai-enterprise-](https://datahack.analyticsvidhya.com/contest/intel-ai-enterprise-meetup-mumbai-invite-only/?utm_source=Sticky_banner1&utm_medium=display&utm_campaign=IntelMum)

[meetup-mumbai-invite-only/?utm\\_source=Sticky\\_banner1&utm\\_medium=display&utm\\_campaign=IntelMum](https://datahack.analyticsvidhya.com/contest/intel-ai-enterprise-meetup-mumbai-invite-only/?utm_source=Sticky_banner1&utm_medium=display&utm_campaign=IntelMum))



([https://courses.analyticsvidhya.com/courses/applied-machine-](https://courses.analyticsvidhya.com/courses/applied-machine-learning-beginner-to-professional?utm_source=Sticky_banner2&utm_medium=display&utm_campaign=Applied_ML)  
science journey.

[learning-beginner-to-professional?utm\\_source=Sticky\\_banner2&utm\\_medium=display&utm\\_campaign=Applied\\_ML](https://courses.analyticsvidhya.com/courses/applied-machine-learning-beginner-to-professional?utm_source=Sticky_banner2&utm_medium=display&utm_campaign=Applied_ML))

## Your Ultimate path for Becoming a DATA Scientist!

Analytics Vidhya has learning path to start your data science journey.

 Download Resource

## ANALYTICS VIDHYA

About Us

(http://www.analyticsvidhya.com/about-me/)

Our Team

(https://www.analyticsvidhya.com/about-me/team/)

Career

(https://www.analyticsvidhya.com/career-analytcs-vidhya/)

Contact Us

(https://www.analyticsvidhya.com/contact/)

Write for us

(https://www.analyticsvidhya.com/about-me/write/)

## DATA SCIENTISTS

Blog

(https://www.analyticsvidhya.com/blog/)

Hackathon

(https://datahack.analyticsvidhya.com/)

Discussions

(https://discuss.analyticsvidhya.com/)

Apply Jobs

(https://www.analyticsvidhya.com/jobs/)

Leaderboard

(https://datahack.analyticsvidhya.com/leaderboard/)

## COMPANIES

Post Jobs

(https://www.analyticsvidhya.com/corporate/)

Trainings

(https://www.analyticsvidhya.com/trainings/)

Hiring Hackathons

(https://www.analyticsvidhya.com/hiring-hackathons/)

Advertising

(https://www.analyticsvidhya.com/contact/)

Reach Us

(https://www.analyticsvidhya.com/contact/)

## JOIN OUR COMMUNITY :

f

(https://www.facebook.com/analyticsvidhya/)

🐦

(https://twitter.com/analyticsvidhya)

(https://www.facebook.com/analyticsvidhya/)

(https://www.facebook.com/analyticsvidhya/)

(https://www.facebook.com/analyticsvidhya/)

Followers

(https://www.facebook.com/analyticsvidhya/)

Followers

(https://www.facebook.com/analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

(https://plus.google.com/+Analyticsvidhya/)

Subscribe to emailer

>

© Copyright 2013-2019 Analytics Vidhya.

Privacy Policy (https://www.analyticsvidhya.com/privacy-policy/)

Terms of Use (https://www.analyticsvidhya.com/terms/)

Refund Policy (https://www.analyticsvidhya.com/refund-policy/)

Don't have an account? Sign up (https://id.a

x

-

(http://play.google.com/store/apps/details?id=com.analyticsvidhya.android)

## Your Ultimate path for Becoming a DATA Scientist!

Download this learning path to start your data science journey.

Name

Email Id

Contact Number



Download Resource